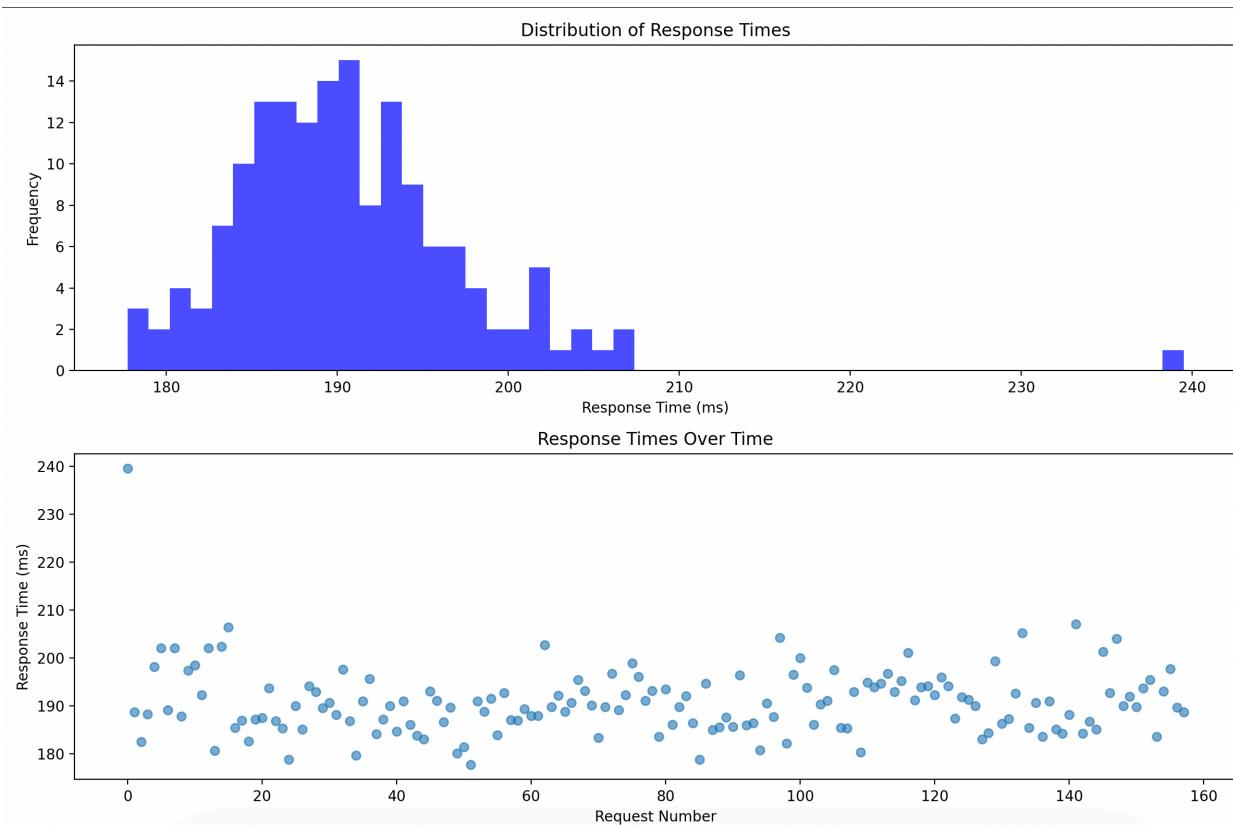


# Building Scalable Distributed Systems- Assignment 1b

## Load Test Results:



## Observations from Load test results:

### 1. Distribution Shape

The histogram shows a mostly normal distribution with a slight right-hand tail. The majority of requests cluster between ~180 ms and ~200 ms, while only a small fraction exceed 200 ms, with one noticeable outlier around ~240 ms.

Only a very small percentage (<5%) of requests can be classified as "slow," indicating generally stable latency with rare spikes rather than persistent slowness.

## **2. Consistency**

The scatter plot shows that response times remain relatively consistent across the entire 30-second window. There is no increasing trend, which suggests the system does not degrade under sustained load. While there are occasional spikes, they are isolated and non-recurring, indicating transient effects rather than systemic performance issues.

## **3. Percentiles**

The median (50th percentile) response time is around ~188–190 ms, while the 95th percentile is closer to ~200–205 ms. The relatively small gap between the median and 95th percentile suggests low variability.

## **4. Infrastructure Impact**

This service is deployed as a single container on a basic EC2 instance, meaning all requests share the same CPU cores, memory, and network interface.

Contributing factors include:

### **1. CPU scheduling and context switching**

- With limited vCPUs, the container competes with:
  - The host OS
  - The container runtime (Docker)
  - Other background system processes
- This can introduce small but measurable latency variations between requests.

### **2. Single-instance execution bottleneck**

- With only one container, requests cannot be distributed.
- Even sequential requests must wait for the same execution context, limiting throughput and increasing sensitivity to transient delays.

## **5. Scaling Implications**

This test uses sequential requests from a single client, which does not reflect real-world traffic patterns. With 100 concurrent users, response times would likely:

- Increase significantly due to CPU contention
- Exhibit a longer tail in the histogram
- Show higher 95th and 99th percentile latencies

Horizontal scaling (multiple containers behind a load balancer) or vertical scaling (larger instance) would be required to maintain current performance levels.

## 6. Network vs Processing

The occasional high-latency outliers suggest a combination of network variability and server-side processing delays.

### **Further investigation could include:**

- Logging server-side request processing time
- Monitoring EC2 CPU and memory usage during load
- Comparing results against a local (same-host) client to isolate network effects